

Learning Work Embeddings

Group No :151

Group Member Names:

Peyala Samarasimha Reddy - 2023AA05072 100% Contribution

Monisha G - 2023AA05536 100% Contribution

Akshay Mohan - 2023AA05315 100% Contribution

Sreelakshmi Ajith - 2023AA05316 100% Contribution

```
In [ ]: import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline

import pandas as pd

import tensorflow as tf
from tensorflow import keras
```

```
In [ ]: tf.random.set_seed(42)
```

Dataset - IMDB

```
In [ ]: imdb = keras.datasets.imdb

max_features = 20000

(Xtrain, Ytrain), (Xtest, Ytest) = imdb.load_data(num_words = max_features)

# Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/i
# 17465344/17464789 [=====] - 8s 0us/step

print(len(Xtrain), len(Ytrain))
print(len(Xtest), len(Ytest))

Xtrain = keras.preprocessing.sequence.pad_sequences(Xtrain, maxlen =25)
Xtest = keras.preprocessing.sequence.pad_sequences(Xtest, maxlen =25)

print(len(Xtrain), len(Ytrain))
print(len(Xtest), len(Ytest))
```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz

17464789/17464789 ————— 0s 0us/step

25000 25000

25000 25000

25000 25000

25000 25000

LSTM

```
In [ ]: lstmModel = keras.models.Sequential()

lstmModel.add(keras.layers.Embedding(input_dim = max_features, output_dim = 128))

lstmModel.add(keras.layers.LSTM(128, dropout=0.2)) #, recurrent_dropout=0.2

lstmModel.add(keras.layers.Dense(1, activation = 'sigmoid'))

lstmModel.summary()
```

Model: "sequential"

Layer (type)	Output Shape	
embedding (Embedding)	?	0 (u
lstm (LSTM)	?	0 (u
dense (Dense)	?	0 (u

Total params: 0 (0.00 B)

Trainable params: 0 (0.00 B)
















Non-trainable params: 0 (0.00 B)

```
In [ ]: # Configure the model for training, by using appropriate optimizers and regularizati
# Available optimizer: adam, rmsprop, adagrad, sgd
# loss: objective that the model will try to minimize.
# Available loss: categorical_crossentropy, binary_crossentropy, mean_squared_error
# metrics: List of metrics to be evaluated by the model during training and testing.

lstmModel.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [ ]: # train the model

history = lstmModel.fit(Xtrain, Ytrain, epochs = 15, batch_size=16, validation_split=
```

Epoch 1/15
1250/1250  **41s** 31ms/step - accuracy: 0.6899 - loss: 0.5659 - val_accuracy: 0.7698 - val_loss: 0.4641
Epoch 2/15
1250/1250  **40s** 30ms/step - accuracy: 0.8353 - loss: 0.3645 - val_accuracy: 0.7626 - val_loss: 0.5003
Epoch 3/15
1250/1250  **37s** 27ms/step - accuracy: 0.9050 - loss: 0.2298 - val_accuracy: 0.7614 - val_loss: 0.6628
Epoch 4/15
1250/1250  **35s** 28ms/step - accuracy: 0.9465 - loss: 0.1436 - val_accuracy: 0.7504 - val_loss: 0.8510
Epoch 5/15
1250/1250  **43s** 29ms/step - accuracy: 0.9631 - loss: 0.0989 - val_accuracy: 0.7372 - val_loss: 1.1171
Epoch 6/15
1250/1250  **37s** 26ms/step - accuracy: 0.9715 - loss: 0.0791 - val_accuracy: 0.7344 - val_loss: 1.0705
Epoch 7/15
1250/1250  **40s** 25ms/step - accuracy: 0.9817 - loss: 0.0528 - val_accuracy: 0.7348 - val_loss: 1.1987
Epoch 8/15
1250/1250  **42s** 26ms/step - accuracy: 0.9878 - loss: 0.0376 - val_accuracy: 0.7370 - val_loss: 1.2957
Epoch 9/15
1250/1250  **41s** 25ms/step - accuracy: 0.9901 - loss: 0.0280 - val_accuracy: 0.7320 - val_loss: 1.3846
Epoch 10/15
1250/1250  **41s** 26ms/step - accuracy: 0.9903 - loss: 0.0254 - val_accuracy: 0.7400 - val_loss: 1.3623
Epoch 11/15
1250/1250  **41s** 26ms/step - accuracy: 0.9928 - loss: 0.0210 - val_accuracy: 0.7314 - val_loss: 1.4050
Epoch 12/15
1250/1250  **43s** 27ms/step - accuracy: 0.9953 - loss: 0.0146 - val_accuracy: 0.7250 - val_loss: 1.6139
Epoch 13/15
1250/1250  **39s** 26ms/step - accuracy: 0.9947 - loss: 0.0148 - val_accuracy: 0.7342 - val_loss: 1.7748
Epoch 14/15
1250/1250  **33s** 26ms/step - accuracy: 0.9960 - loss: 0.0134 - val_accuracy: 0.7344 - val_loss: 1.7806
Epoch 15/15
1250/1250  **41s** 26ms/step - accuracy: 0.9972 - loss: 0.0093 - val_accuracy: 0.7352 - val_loss: 1.8421



In []: *# plotting training and validation loss*

```
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, color='red', label='Training loss')
plt.plot(epochs, val_loss, color='green', label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

plotting training and validation Accuracy

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, acc, color='red', label='Training Accuracy')
```

```
plt.plot(epochs, val_acc, color='green', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



```
In [ ]: testLoss, testAccuracy = lstmModel.evaluate(Xtest, Ytest)
print(testLoss, testAccuracy)
```

782/782 ————— 6s 8ms/step - accuracy: 0.7313 - loss: 1.8497
1.821761131286621 0.7311199903488159

MODIFICATIONS

Use a different optimizer

```
In [ ]: #ADAM replaced by SGD
from tensorflow.keras.optimizers import SGD

lstmModel.compile(optimizer=SGD(learning_rate=0.01), loss='binary_crossentropy', metr
```

```
In [ ]: #RMSProp
from tensorflow.keras.optimizers import RMSprop
lstmModel.compile(optimizer=RMSprop(learning_rate=0.001), loss='binary_crossentropy',
```

Train for more epochs


```
In [ ]: #epoch increased to 20
history_modified = lstmModel.fit(Xtrain, Ytrain, epochs = 20, batch_size=64, validati
```

```
Epoch 1/20
113/113 ————— 15s 45ms/step - accuracy: 0.9934 - loss: 0.0222 - val_acc
uracy: 0.7368 - val_loss: 1.7653
Epoch 2/20
113/113 ————— 21s 47ms/step - accuracy: 0.9938 - loss: 0.0186 - val_acc
uracy: 0.7374 - val_loss: 1.7368
Epoch 3/20
113/113 ————— 20s 45ms/step - accuracy: 0.9952 - loss: 0.0137 - val_acc
uracy: 0.7386 - val_loss: 1.7177
Epoch 4/20
113/113 ————— 22s 49ms/step - accuracy: 0.9955 - loss: 0.0130 - val_acc
uracy: 0.7382 - val_loss: 1.7081
Epoch 5/20
113/113 ————— 19s 44ms/step - accuracy: 0.9965 - loss: 0.0123 - val_acc
uracy: 0.7380 - val_loss: 1.7071
Epoch 6/20
113/113 ————— 14s 44ms/step - accuracy: 0.9963 - loss: 0.0102 - val_acc
uracy: 0.7370 - val_loss: 1.7037
Epoch 7/20
113/113 ————— 20s 43ms/step - accuracy: 0.9973 - loss: 0.0090 - val_acc
uracy: 0.7378 - val_loss: 1.7061
Epoch 8/20
113/113 ————— 21s 44ms/step - accuracy: 0.9964 - loss: 0.0105 - val_acc
uracy: 0.7384 - val_loss: 1.7058
Epoch 9/20
113/113 ————— 20s 44ms/step - accuracy: 0.9969 - loss: 0.0090 - val_acc
uracy: 0.7388 - val_loss: 1.7076
Epoch 10/20
113/113 ————— 20s 43ms/step - accuracy: 0.9964 - loss: 0.0095 - val_acc
uracy: 0.7388 - val_loss: 1.7082
Epoch 11/20
113/113 ————— 20s 43ms/step - accuracy: 0.9973 - loss: 0.0081 - val_acc
uracy: 0.7396 - val_loss: 1.7095
Epoch 12/20
113/113 ————— 14s 43ms/step - accuracy: 0.9977 - loss: 0.0068 - val_acc
uracy: 0.7402 - val_loss: 1.7122
Epoch 13/20
113/113 ————— 20s 43ms/step - accuracy: 0.9969 - loss: 0.0087 - val_acc
uracy: 0.7396 - val_loss: 1.7199
Epoch 14/20
113/113 ————— 21s 43ms/step - accuracy: 0.9974 - loss: 0.0085 - val_acc
uracy: 0.7402 - val_loss: 1.7223
Epoch 15/20
113/113 ————— 21s 45ms/step - accuracy: 0.9974 - loss: 0.0084 - val_acc
uracy: 0.7404 - val_loss: 1.7233
Epoch 16/20
113/113 ————— 14s 43ms/step - accuracy: 0.9971 - loss: 0.0089 - val_acc
uracy: 0.7408 - val_loss: 1.7275
Epoch 17/20
113/113 ————— 14s 43ms/step - accuracy: 0.9980 - loss: 0.0061 - val_acc
uracy: 0.7408 - val_loss: 1.7312
Epoch 18/20
113/113 ————— 21s 44ms/step - accuracy: 0.9977 - loss: 0.0066 - val_acc
uracy: 0.7402 - val_loss: 1.7309
Epoch 19/20
113/113 ————— 20s 43ms/step - accuracy: 0.9972 - loss: 0.0074 - val_acc
uracy: 0.7398 - val_loss: 1.7365
Epoch 20/20
113/113 ————— 14s 46ms/step - accuracy: 0.9976 - loss: 0.0073 - val_acc
uracy: 0.7396 - val_loss: 1.7355
```


```
In [ ]: #Rmsprop
```

```
history_rms = lstmModel.fit(Xtrain, Ytrain, epochs=20, batch_size=64, validation_split=0.1)
```


Epoch 1/20

313/313  **16s** 48ms/step - accuracy: 0.9987 - loss: 0.0044 - val_accuracy: 0.7362 - val_loss: 2.0575


Epoch 2/20

313/313  **15s** 47ms/step - accuracy: 0.9994 - loss: 0.0021 - val_accuracy: 0.7344 - val_loss: 2.1209


Epoch 3/20

313/313  **20s** 47ms/step - accuracy: 0.9998 - loss: 0.0013 - val_accuracy: 0.7382 - val_loss: 2.1908


Epoch 4/20

313/313  **21s** 48ms/step - accuracy: 0.9999 - loss: 6.4327e-04 - val_accuracy: 0.7414 - val_loss: 2.2825


Epoch 5/20

313/313  **21s** 51ms/step - accuracy: 0.9999 - loss: 6.0977e-04 - val_accuracy: 0.7394 - val_loss: 2.3280


Epoch 6/20

313/313  **19s** 47ms/step - accuracy: 0.9994 - loss: 0.0012 - val_accuracy: 0.7416 - val_loss: 2.2973


Epoch 7/20

313/313  **21s** 48ms/step - accuracy: 0.9998 - loss: 6.7363e-04 - val_accuracy: 0.7410 - val_loss: 2.3313


Epoch 8/20

313/313  **21s** 50ms/step - accuracy: 0.9999 - loss: 4.0915e-04 - val_accuracy: 0.7414 - val_loss: 2.4033


Epoch 9/20

313/313  **20s** 47ms/step - accuracy: 0.9999 - loss: 5.3054e-04 - val_accuracy: 0.7416 - val_loss: 2.4885


Epoch 10/20

313/313  **20s** 47ms/step - accuracy: 0.9999 - loss: 3.6763e-04 - val_accuracy: 0.7404 - val_loss: 2.4748


Epoch 11/20

313/313  **21s** 47ms/step - accuracy: 0.9999 - loss: 2.6144e-04 - val_accuracy: 0.7408 - val_loss: 2.4864


Epoch 12/20

313/313  **15s** 47ms/step - accuracy: 0.9999 - loss: 3.9519e-04 - val_accuracy: 0.7404 - val_loss: 2.5714


Epoch 13/20

313/313  **21s** 48ms/step - accuracy: 0.9999 - loss: 2.6912e-04 - val_accuracy: 0.7396 - val_loss: 2.5880


Epoch 14/20

313/313  **21s** 49ms/step - accuracy: 0.9999 - loss: 3.6115e-04 - val_accuracy: 0.7420 - val_loss: 2.5947


Epoch 15/20

313/313  **20s** 48ms/step - accuracy: 0.9999 - loss: 1.6911e-04 - val_accuracy: 0.7364 - val_loss: 2.6168


Epoch 16/20

313/313  **15s** 49ms/step - accuracy: 0.9995 - loss: 0.0026 - val_accuracy: 0.7392 - val_loss: 2.5877


Epoch 17/20

313/313  **15s** 47ms/step - accuracy: 0.9999 - loss: 1.9724e-04 - val_accuracy: 0.7390 - val_loss: 2.6479


Epoch 18/20

313/313  **21s** 50ms/step - accuracy: 0.9999 - loss: 1.1133e-04 - val_accuracy: 0.7384 - val_loss: 2.7109

Epoch 19/20

313/313  **20s** 47ms/step - accuracy: 0.9999 - loss: 1.8545e-04 - val_accuracy: 0.7398 - val_loss: 2.7085

Epoch 20/20

313/313  **21s** 47ms/step - accuracy: 0.9999 - loss: 1.3149e-04 - val_accuracy: 0.7404 - val_loss: 2.7104

Use the learned network to predict the sentiments for new sentences

```
In [ ]: # Example new sentences
new_sentences = ["The movie was fantastic!", "I am feeling sad."]

# Preprocess the new sentences (tokenization, padding, etc.)
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.preprocessing.text import Tokenizer

tokenizer = Tokenizer(num_words=10000)

new_sequences = tokenizer.texts_to_sequences(new_sentences)
new_padded_sequences = pad_sequences(new_sequences, maxlen=100)

# Use the model to predict sentiments
predictions = lstmModel.predict(new_padded_sequences)

# Output predictions
for i, sentence in enumerate(new_sentences):
    print(f"Sentence: {sentence}")
    print(f"Predicted Sentiment: {'Positive' if predictions[i] > 0.5 else 'Negative'}
```

1/1 ————— 0s 165ms/step

Sentence: The movie was fantastic!

Predicted Sentiment: Positive

Sentence: I am feeling sad.

Predicted Sentiment: Positive

Graph plot after modification


```
In [ ]: # plotting training and validation loss

loss = history_rms.history['loss']
val_loss = history_rms.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, color='red', label='Training loss')
plt.plot(epochs, val_loss, color='green', label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()

# plotting training and validation Accuracy

acc = history_rms.history['accuracy']
val_acc = history_rms.history['val_accuracy']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, acc, color='red', label='Training Accuracy')
plt.plot(epochs, val_acc, color='green', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

Training and Validation loss



Training and Validation Accuracy

